



**Institute** of  
**Data**

---

2023



# Data Science and AI

## Module 10

---

### Machine Learning Deployment and Cloud Computing

---



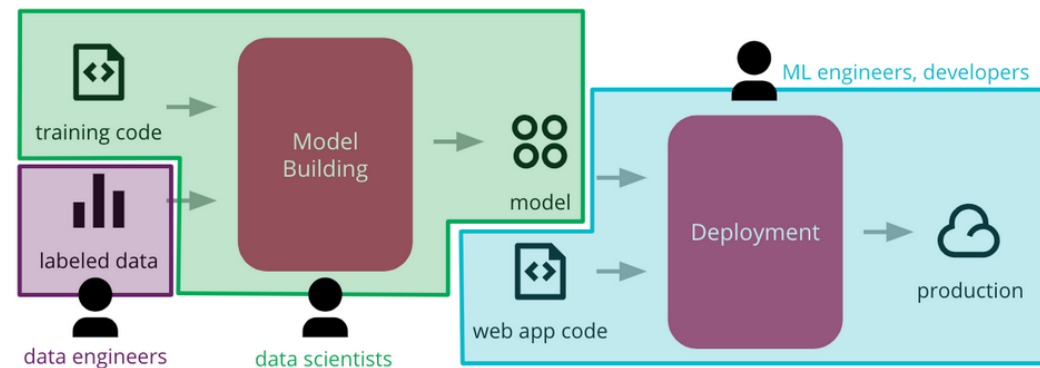

# Agenda: Module 10 – Machine Learning Deployment and Cloud Computing

- Introduction
- Cloud computing
- Deployment
- DevOps, AutoML
- Cluster Computing
- Streaming Data



# What is Machine Learning Deployment?

The process of making machine learning models and analyses easier to use and more accessible, bridging the gap between models and consumers



<https://martinfowler.com/articles/cd4ml.html>



# Key Machine Learning Engineering Skills

- Well-versed in Cloud Computing and Big Data technologies such as Hadoop or Spark
- Machine learning – knowledge of algorithms and how to code them using libraries (e.g. sklearn)
- Data engineering – automating the processes of feeding data in and feature engineering/selection
- Software engineering
  - scaling out models into a “production environment” where it is in operation for consumption by their intended users
  - optimising code for performance



# Cloud Computing

- Basics
- Virtualisation
- Cloud Computing Service Models
  - SaaS, PaaS, IaaS
- Introduction to AWS
  - S3
  - Machine Learning Services



# The Challenges of Real-World Data and Computation

- “Big” Data is more than volume – it is about its complexity making it untenable to process using traditional means
  - **Volume** (requires multiple machines to store and process)
  - **Velocity** (data in motion)
  - **Variety** (often unstructured)
  - **Veracity** (often unreliable or incomplete)
- Challenge for data science – making models work in a dynamic real-time environment
  - Changes to data, algorithms, code
- Utilising a **cloud computing** environment addresses some of these challenges



# Cloud Computing Basics

- **Cloud Computing** – the delivery of IT services (e.g. software, storage, computing processors) over a computer network (usually the Internet or Intranet)
- Characteristics (as per NIST – National Institute of Standards and Technology)
  - on-demand self-service provisioning of resources
  - broad network access
  - resource pooling (grouping together resources to improve the quality of service)
  - rapid elasticity (ability to scale up or down easily)
  - measured service (usage of resources is calculated and billed accordingly)
- Typically makes use of **virtualisation**





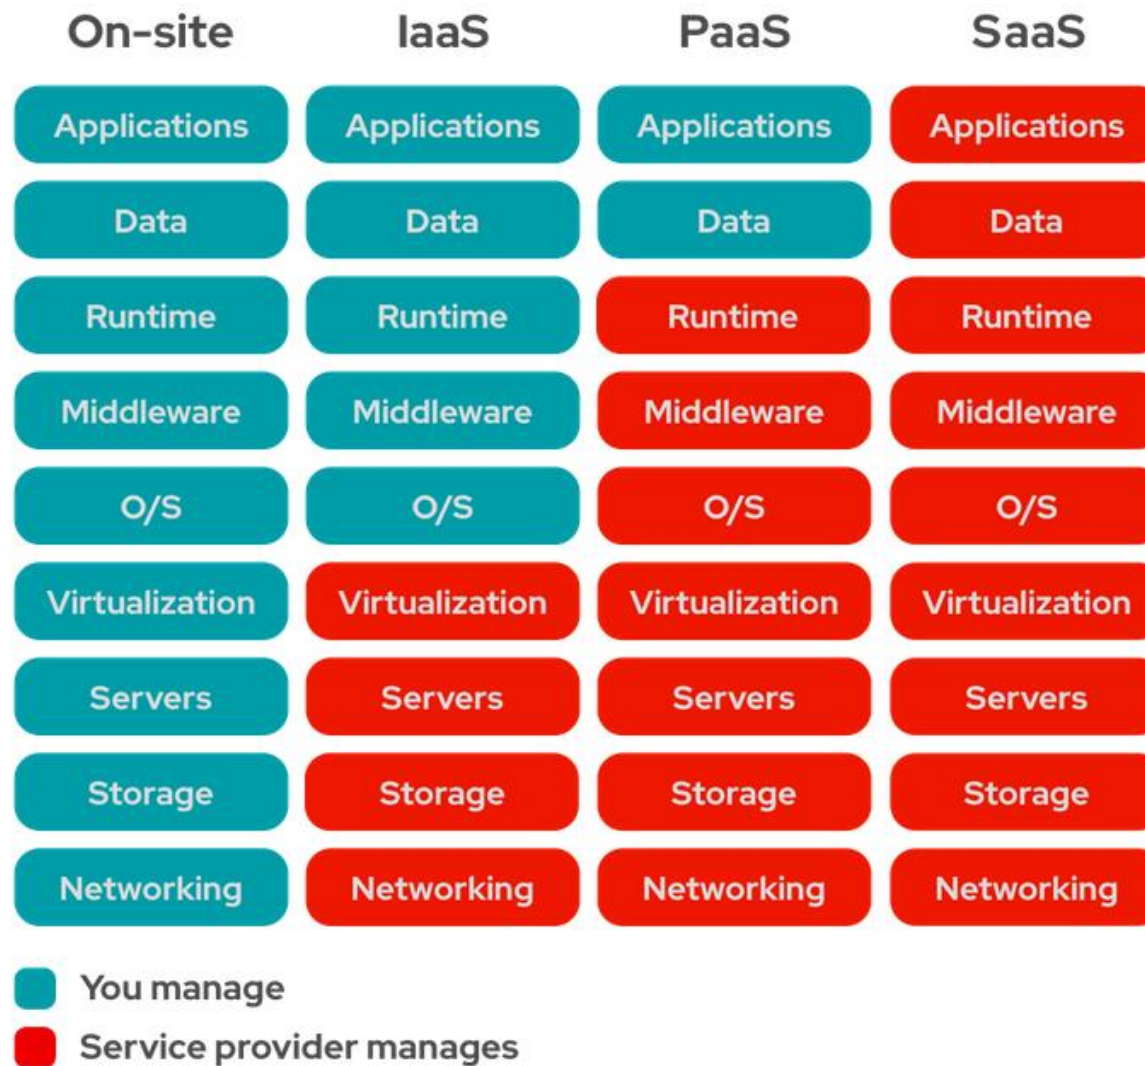
# Virtualisation

- **Virtualisation** allows multiple operating systems to be run on the same physical hardware of a computer
- A **virtual machine** (VM) is a software program or operating system running on hardware that behaves as though it were a single computer that runs its own applications
  - Multiple virtual machine **instances** may exist on the same physical host machine
  - If an application crashes on one VM it does not affect other applications on different VMs
- Allows for a more efficient use of computing resources



# Cloud Computing Service Models

- **SaaS** – entire application made available over a web browser (e.g. Gmail, Office 365)
- **PaaS** – software development and deployment (e.g. AWS Elastic Beanstalk, Heroku, Render, Google App Engine)
- **IaaS** – provides storage, networking, virtualised hardware (e.g. AWS, Microsoft Azure, Google Compute Engine)



<https://www.redhat.com/en/topics/cloud-computing/iaas-vs-paas-vs-saas>



# Instance Specifications

- Number and Type of processor (CPU vs GPU vs TPU)
- Memory (RAM or Storage)
- Type of Operating System (Linux/Unix, MacOS, Windows)
- Bandwidth
- Instances optimised for Compute, Memory (large datasets in memory) or Storage
- AWS Example: <https://aws.amazon.com/ec2/instance-types/>




# Leading Cloud Computing providers

File storage

Computation

List of services



<https://aws.amazon.com/products/>



File



Queue



Table



Blob

Storage Types



Azure VM

<https://azure.microsoft.com/en-au/services/>



Google Cloud Storage



Google Compute Engine

<https://cloud.google.com/products>



# AWS S3 Buckets

- Simple Storage Service – does not provide computation
- **Buckets** store objects (data) and associated metadata
- Files are referenced as keys – there is no hierarchical structure
  - e.g. though dir/f1 and dir/f2 are allowed filenames, they do not share a common folder dir
- Its name should be globally unique
- Pricing: <https://aws.amazon.com/s3/pricing/>



# Amazon Web Services



## Machine Learning

Build with powerful services and platforms, and the broadest machine learning framework support anywhere.

[Learn More](#)



## Analytics & Data Lakes

Securely store, categorize, and analyze all your data in one, centralized repository.

[Learn More](#)



## Internet of Things

A system of ubiquitous devices connecting the physical world to the cloud.

[Learn More](#)



## Serverless Computing

Build and run applications and services without thinking about servers.

[Learn More](#)



## Containers

Package and deploy applications that are lightweight and provide a consistent, portable software environment for applications to easily run and scale anywhere.

[Learn More](#)



## Enterprise Applications

Build with a mature set of services specifically designed for the unique security, compliance, privacy, and governance requirements of large organizations.

[Learn More](#)



## Storage

Durable, cost-effective options for backup, disaster recovery, and data archiving at petabyte scale.

[Learn More](#)



## Windows Workloads

Flexible, scalable compute capacity for Microsoft applications. Easily manage and secure Windows workloads.

[Learn More](#)



# AWS for Machine Learning

- **SageMaker** – notebooks (Jupyter and JupyterLab)
- **Amazon Rekognition** – image and video analysis such as facial recognition
- **Amazon Comprehend** – NLP
- **Amazon Textract** – text extraction from images or pdf files
- Amazon Lex – Building conversational interfaces (e.g. chatbots) into an application using voice and text
- Amazon Polly – text to speech
- Amazon Transcribe – speech to text
- AWS DeepLens – Deep-Learning-enabled video camera
- Amazon Personalize – individual recommendations
- Amazon Forecast – time series forecasting



# The AWS ML Stack

## AI SERVICES

VISION	SPEECH		TEXT		SEARCH <small>NEW</small>	CHATBOTS	PERSONALIZATION	FORECASTING	FRAUD <small>NEW</small>	DEVELOPMENT <small>NEW</small>	CONTACT CENTERS	
												
Amazon Rekognition	Amazon Polly	Amazon Transcribe <small>+Medical</small>	Amazon Comprehend <small>+Medical</small>	Amazon Translate	Amazon Textract	Amazon Kendra	Amazon Lex	Amazon Personalize	Amazon Forecast	Amazon Fraud Detector	Amazon CodeGuru	Amazon Connect <small>with Contact Lens</small>

## ML SERVICES

 Amazon SageMaker	Ground Truth data labelling	ML Marketplace	SageMaker Studio IDE <span>NEW</span>							SageMaker Neo
			Built-in algorithms	SageMaker Notebooks <span>NEW</span>	SageMaker Experiments <span>NEW</span>	Model tuning	SageMaker Autopilot <span>NEW</span>	Model hosting	SageMaker Model Monitor <span>NEW</span>	

## ML FRAMEWORKS & INFRASTRUCTURE

 TensorFlow <small>NEW</small>  mxnet <small>NEW</small>  PYTORCH	 GLUON  Keras	Deep Learning AMIs & Containers	GPUs & CPUs	Elastic Inference	Inferentia	FPGA
--	--	---------------------------------	-------------	-------------------	------------	------

<https://www.forbes.com/sites/moorinsights/2020/03/04/surprise-aws-leads-in-cloud-ai-services-ranking/>

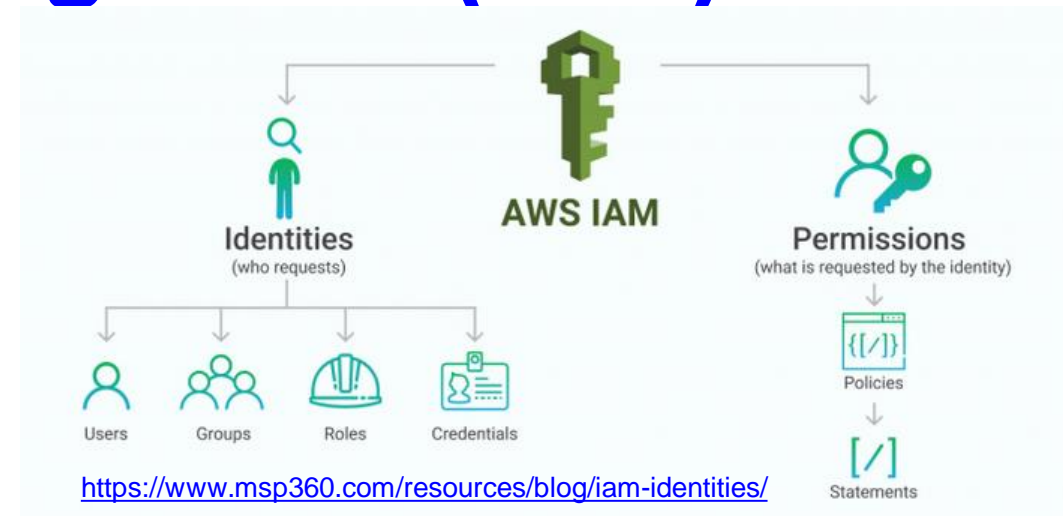
© 2023 Institute of Data





# Identity and Access Management (IAM)

- Access in AWS is managed by creating policies and attaching them to IAM identities (users, groups of users, or roles) or AWS resources.
- **User** – an individual
- **Group** – a set of users sharing common roles and policies
- **Role** – a set of permissions for granting access to resources – not associated with any specific user or group and less permanent
- **Policy** – a document listing permissions (e.g. creating or updating a database table)



Account ID: 123456789012

## Identity-based policies

<b>John Smith</b> Can List, Read On Resource X
<b>Carlos Salazar</b> Can List, Read On Resource Y,Z
<b>MaryMajor</b> Can List, Read, Write On Resource X,Y,Z
<b>ZhangWei</b> No policy

## Resource-based policies

<b>Resource X</b> JohnSmith: Can List, Read MaryMajor: Can List, Read
<b>Resource Y</b> CarlosSalazar: Can List, Write ZhangWei: Can List, Read
<b>Resource Z</b> CarlosSalazar: Denied access ZhangWei: Allowed full access



# Lab 10.1: AWS SageMaker and Managed Services

## Purpose

- Gain familiarity with AWS SageMaker and some Machine Learning Managed Services via the Console

## Resources

- COCO Dataset (images)

## Materials

- Jupyter Notebook (Lab-10\_1)
- Sample image files



# Deployment

- Challenges
- Considerations
- Pipelining
- Small-scale deployment via a REST API
- DevOps and CI/CD
- Monitoring



# Deployment

- **Deployment** refers to making models available in production environments, where they can provide exposure to more users or be consumed by other software systems.
- Only after this deployment do models add business value.
- Types of deployment depend on the frequency of model training:
- **Batch** – large amount of data is required for a prediction (e.g. a user's entire transaction history), latency does not matter
- **Real-time** – a small amount of data is required (e.g. choosing an online ad to show to a customer)



# Challenges of Deployment

- Requires coordination between data scientists, IT teams, software developers, and business professionals to ensure a model works reliably in the organisation's production environment
- Sometimes a model is written in one language (e.g. R, Python) but is deployed in another (e.g. Java, C++)
- Incompatibility of code across different machines (operating systems, versions of libraries etc)
  - -> Containerisation



# Challenges of Deployment

- Processing power (e.g. GPUs, TPUs) might be scarce and expensive for deployment at a large scale
- Scaling to accommodate a large number of users
- Data leakage or contamination – information about a target variable is present when it should not (e.g. information from the future, features depend on data from the holdout set)
- Reproducibility of results



# Deployment Considerations

- Will the deployment deliver business value? Success metrics?
- What software tools to use?
- What are the dependencies (for code and data)?
- How will the results be consumed?
  - A dashboard?
  - Linked to an alert or triggering another software process?



# Deployment Considerations

- Hosted on the cloud or on-premise?
- How to scale it to more users?
- How to guarantee security?
- How often should the model be retrained?
- How to monitor performance?





# Model pipelining

- To ensure reproducibility avoid transforming data manually (e.g. find/replace operations, editing a Spreadsheet in Excel)
- Make use of pipelines consisting of steps such as the following
  - Read in the data (possibly using a feature store)
  - Merge multiple sources
  - Completing missing values
  - Address class imbalance (e.g. resampling)
  - Feature selection/engineering
  - Dimension Reduction
  - Model training with optimised hyperparameters
- In Python: `from sklearn.pipeline import Pipeline`



# Small-Scale Deployment

- Easiest way is to create a web REST API endpoint using a Web framework such as Flask
  - Enables real-time predictions but would not scale up to a large number of users

Basic steps:

- Set up a reproducible pipeline from data to prediction
- Save the model pipeline in binary format (pickle or joblib)
- Build an image (container) with model and library dependencies
- Have the image hosted on a Platform as a Service (e.g. Heroku, Render, AWS Elastic Beanstalk)



# Model as Data – what needs to be stored?

- Persistence – saving a model for future use without having to retrain
- Linear regression – a list of coefficients
  - Predict by applying a linear combination of the inputs
- KNN – the entire dataset
  - Predict by finding the k nearest neighbours to a given input
- Decision tree – the set of decisions at each node
  - Predict by traversing the tree according to the outcome of each decision
- Neural network – all weights and activation functions between neurons
  - Predict by applying the linear combinations and activation functions to a given input




# Containers

- A container is a set of software processes isolated from the rest of an operating system
- It contains all the files necessary to support the processes (model + library dependencies)
- Different from a virtual machine in that they do not require a full operating system – a single operating system may run separate container instances
- Programs inside a container only have visibility to contents of the container and devices connected to it



# Containers

- Advantage: avoiding environment-related issues - if containerised code works on one machine, it will run on another irrespective of the characteristics of the machine
- Docker is the most popular containerization service
- Docker Tutorial: [docker-curriculum.com](https://docker-curriculum.com)





# Deployment via Flask

- Flask is a web framework allowing one to create small-scale web apps in Python
- Allows one to create a REST API service that is deployed on a web server
- For larger applications one can use Gunicorn with Nginx as the web server

```
@app.route('/')  
def home():  
    return render_template('index.html')
```

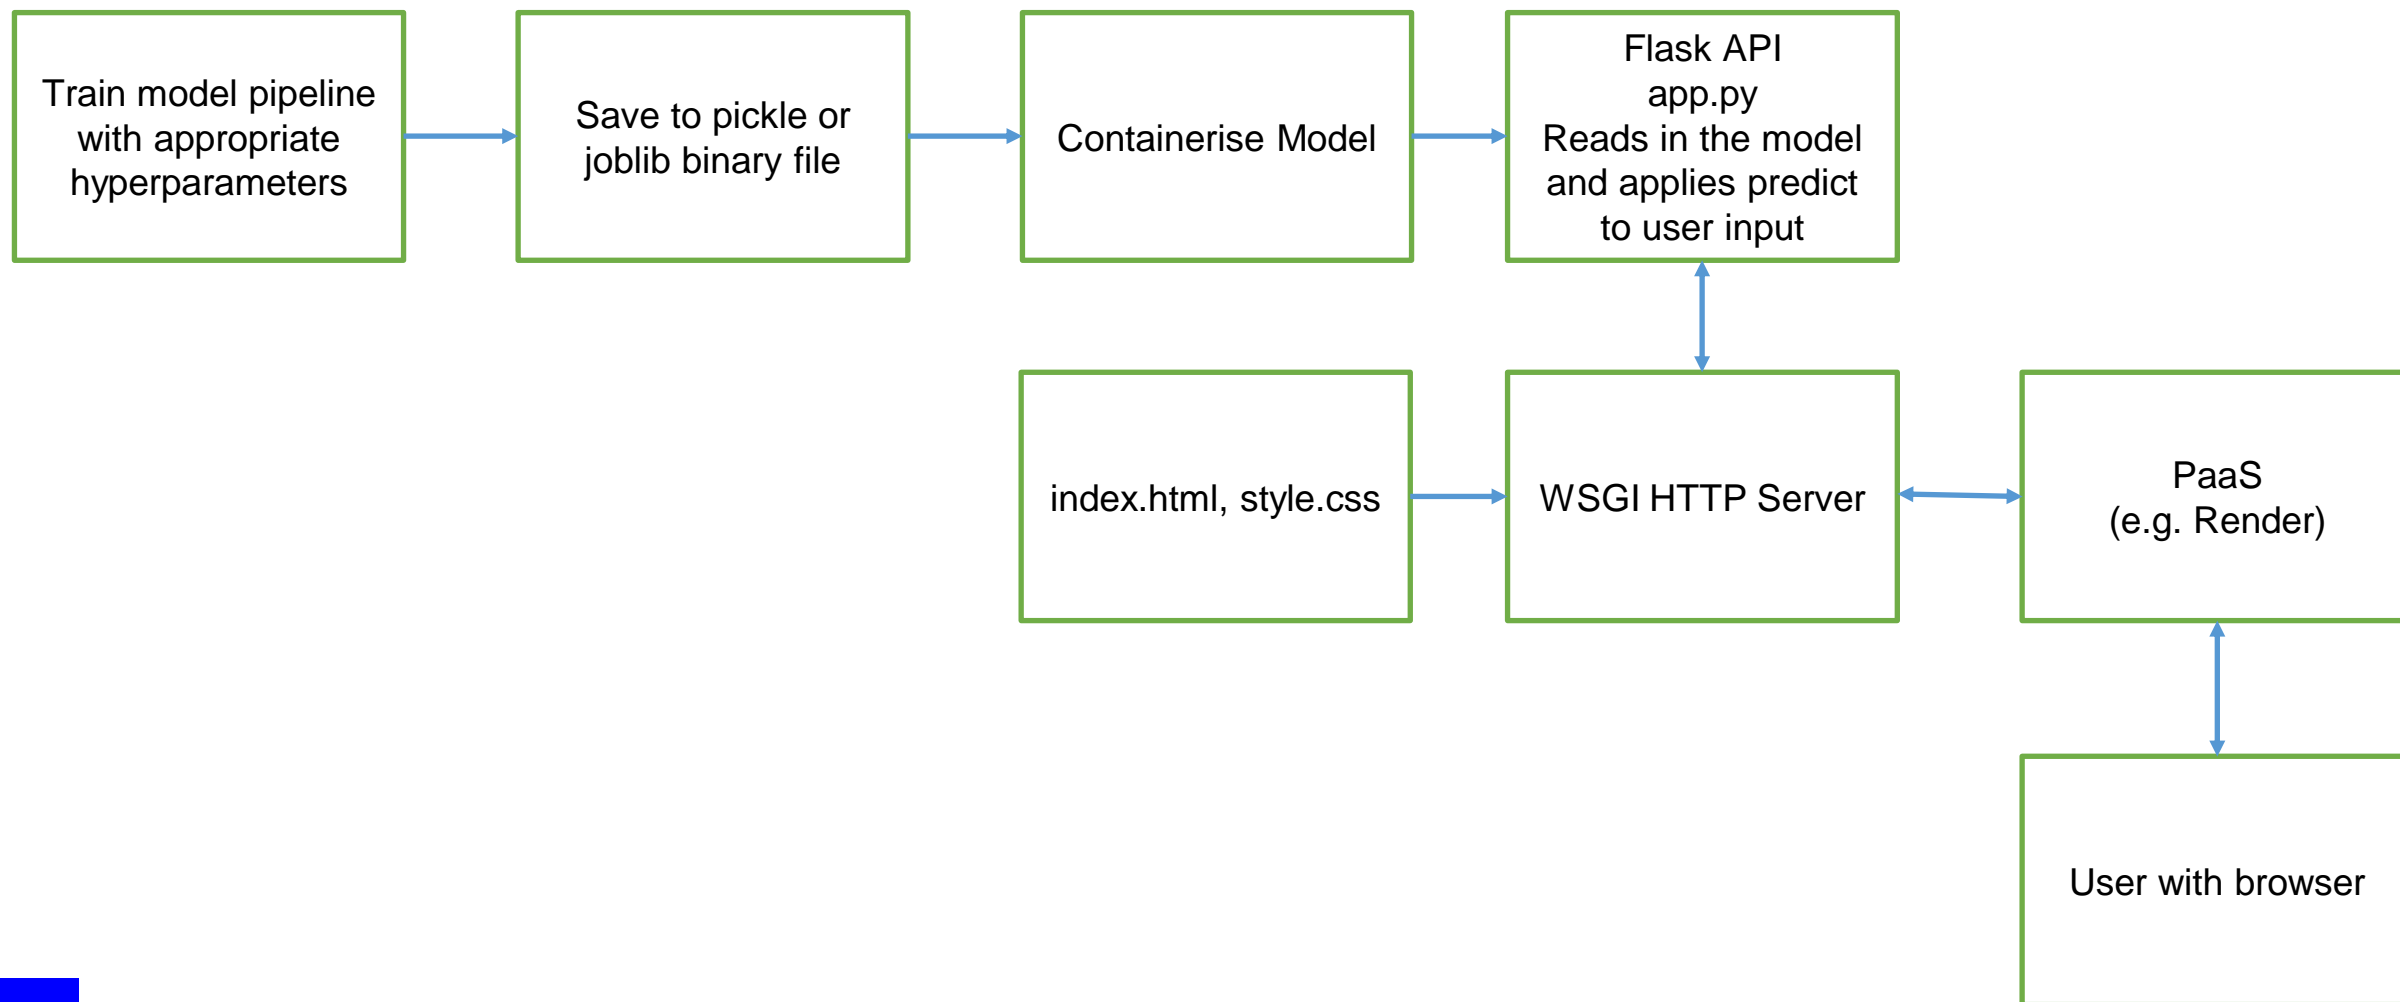
```
@app.route('/predict', methods=['POST'])  
def predict():  
    ...
```

@app.route  
allows Flask to  
assign functions  
to URLs





# Deployment via Flask





# Model Monitoring and Maintenance

- Monitoring and scoring models while in production ensures they are behaving as intended
- Model/data drift – the statistics of input data no longer match the training data – hence the need to retrain
  - Statistical tests can check for the closeness between distributions of training data and data observed in production
- Concept drift – change in the meaning/interpretation of the target variable over time
  - e.g. in detecting rare events the definition of an anomaly changes over time
- Usually the remedy is more frequent training but on occasions the model needs to be redesigned



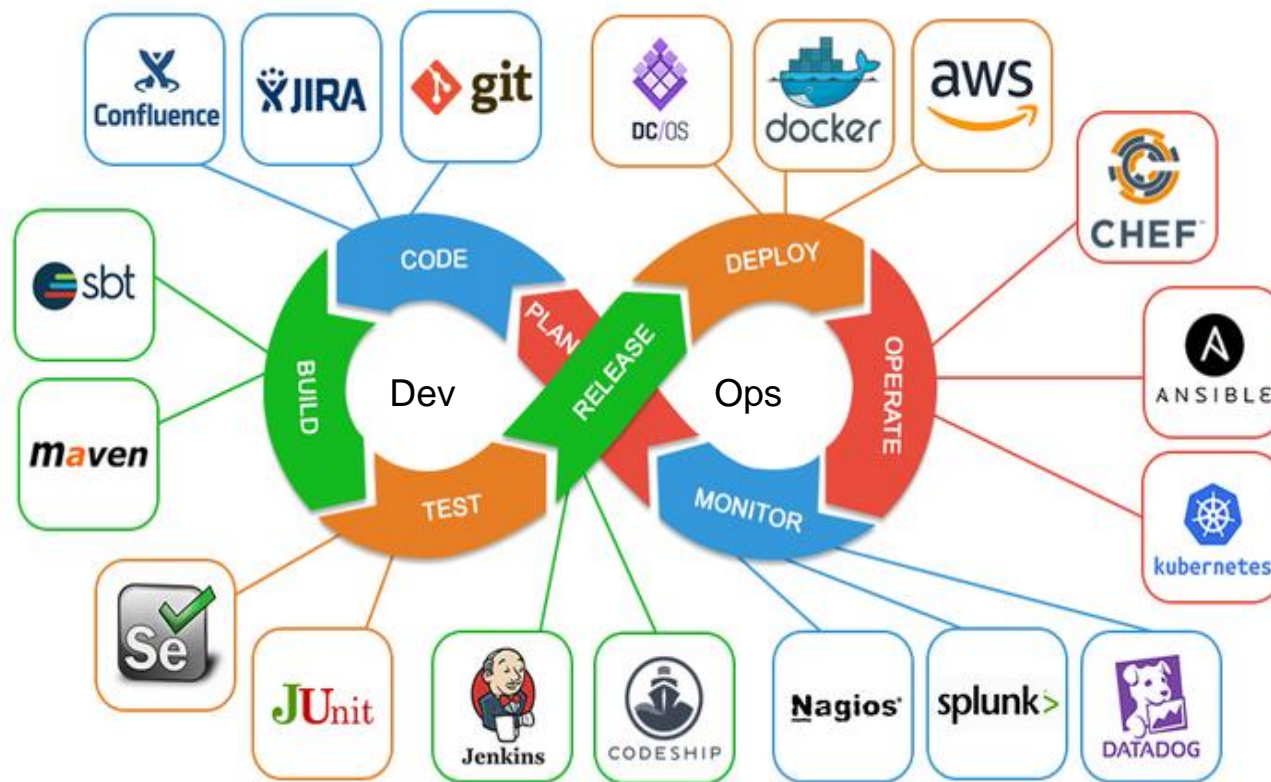


# Larger scale: DevOps and CI/CD

- DevOps is a combination of Software Development (Dev) and IT Operations (Ops)
- IT Operations – services provided by an IT department for its customers
- DevOps aims to shorten the systems development life cycle and provide continuous delivery while maintaining high software quality
- CI/CD (Continuous Integration/Continuous Deployment) automates the process of building-testing-deploying (e.g in Dev/Test/Prod environments), shortening the life cycle of software development while maintaining quality
- New ideas can be deployed in production more rapidly, providing value to the user



# The DevOps Lifecycle with Sample Tools




<https://dzone.com/articles/why-should-i-learn-devops>

See also: <https://i.redd.it/92ja1d3v0x141.jpg>




# Orchestration

- A web service can scale by orchestration – a cluster of containers can be created on demand based on a copy of the original image container
- This also ensures fault tolerance should one container fail
- Kubernetes and Amazon ECS (Elastic Container Service) are well known orchestration systems



1. Create a Kubernetes cluster




2. Deploy an app




3. Explore your app



4. Expose your app publicly



5. Scale up your app



6. Update your app

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>



# Enabling Security

Ways to ensure secure model deployment include

- Cryptographically encrypting the model
- Hosting with a secure protocol (https vs http)
- Providing access control (e.g. token-based API)
- Including password protections



# Machine Learning Platforms

Developed internally

- Google: [TFX](#) (TensorFlow Extended)
- Uber: [Michelangelo](#)
- Databricks: [MLFlow](#)
- Facebook: [FBLearner Flow](#)

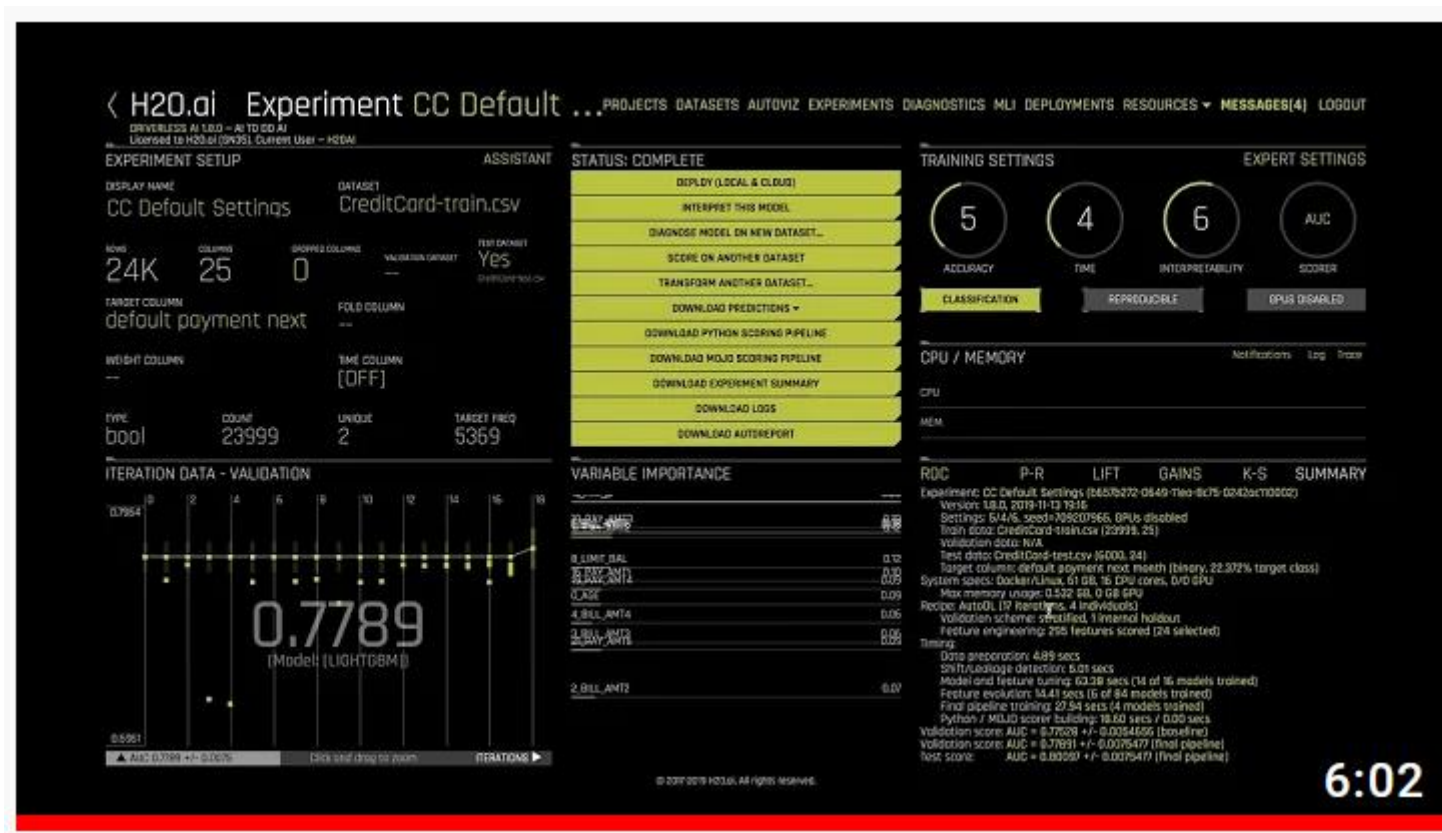


# AutoML

- Tools exist to automate the process from data ingestion to deployment
    - Most useful for data clean-up and hyperparameter tuning
  - Data scientists are still needed for problem definition, appropriate feature engineering requiring domain knowledge
- 
- Google Cloud AutoML
  - Microsoft Azure AutoML
  - AWS SageMaker Autopilot
  - DataRobot
  - H2O.ai
  - TPOT
  - Auto-Sklearn



# One example of AutoML - H2O.ai demo



<https://www.youtube.com/watch?v=ZqCoFp3-rGc>



# Lab 10.2: Model Deployment

- Purpose
  - Learn how to deploy a simple app on a cloud platform
- Resources
  - Reviews Data – sentiments.csv
- Materials
  - Jupyter Notebook (Lab 10\_2)
  - Several Python scripts and text files





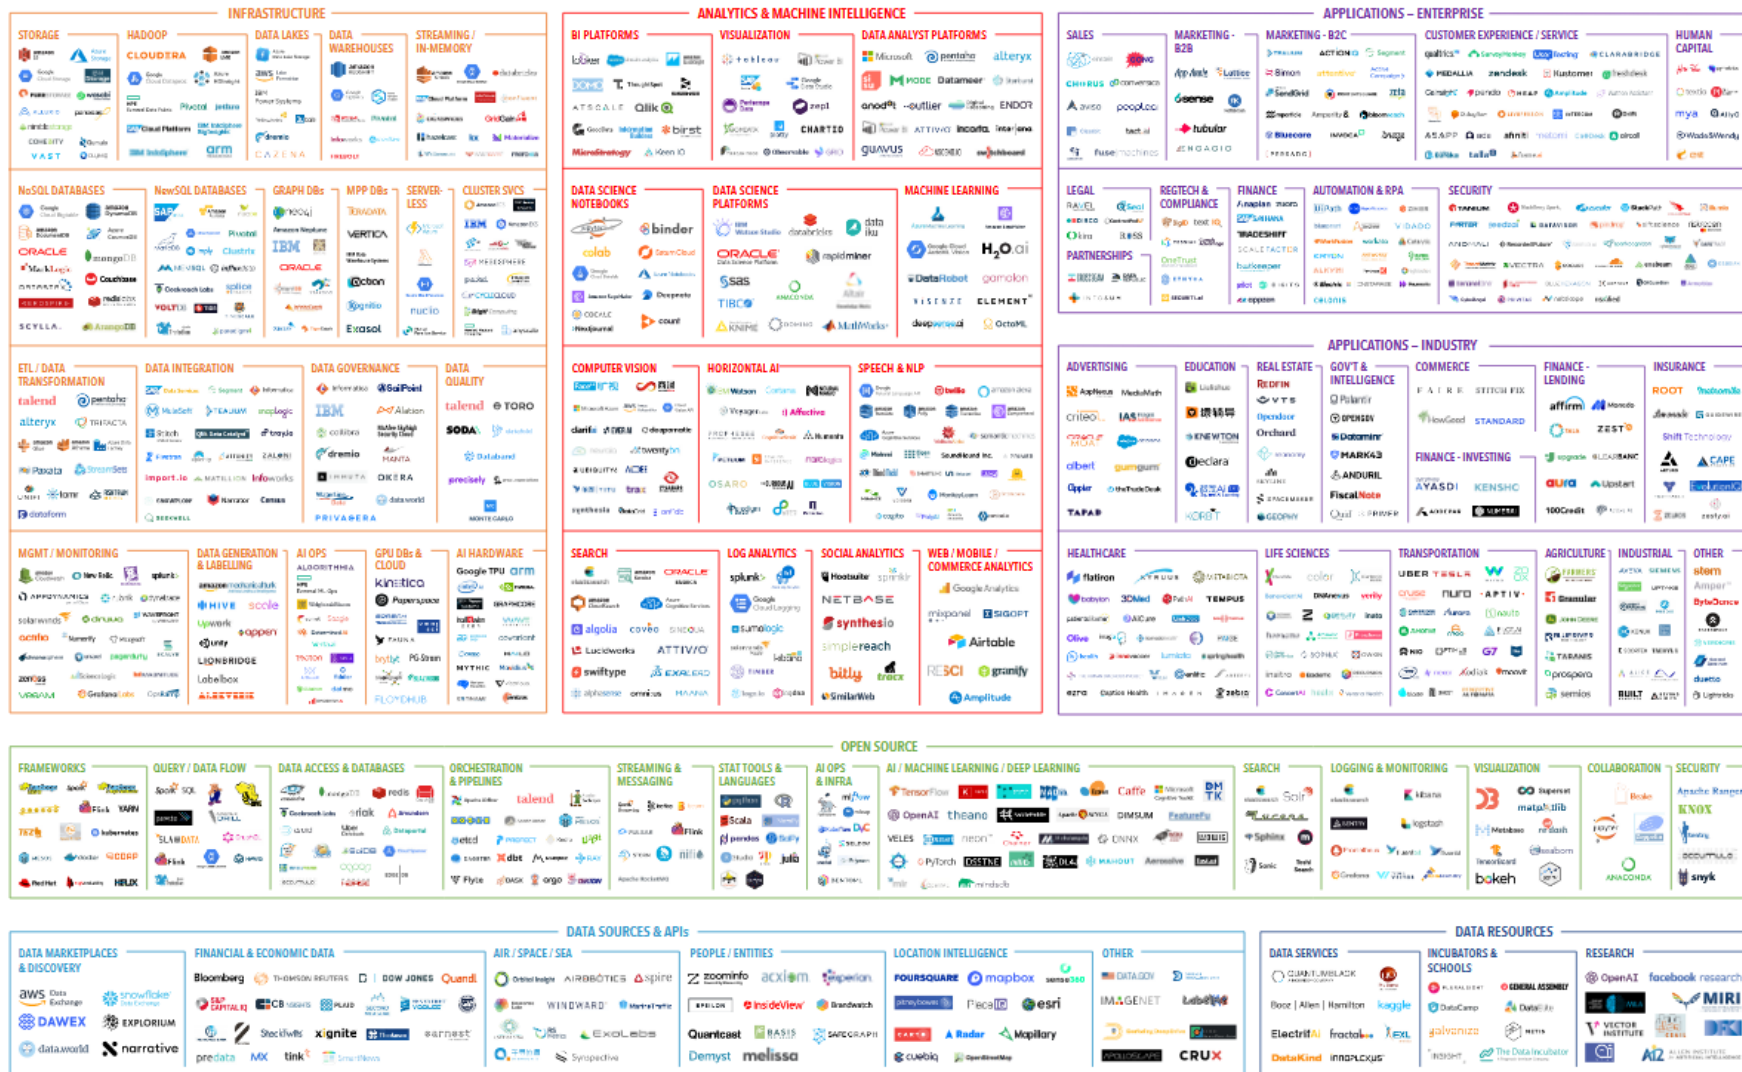
# Cluster Computing

- Introduction
- Hadoop
- Spark
- PySpark
- Amazon EMR



# The Data & AI Landscape is vast!

DATA & AI LANDSCAPE 2020



<https://mattturck.com/data2020/>



# Introduction

- **Cluster computing** – two or more computing nodes (servers) connected via a Fast Area Network, usually performing the same task
- Chief benefits – speed, scalability, flexibility over traditional mainframe solutions
- A **load balancer** may be used to spread the computing load evenly across the cluster for better overall performance
- Applications in Machine Learning
  - Working with large datasets
  - Deep learning (often nodes will have GPUs or TPUs)



# Parallel Computation

- MapReduce: splits a large task into multiple map and reduce tasks that can be distributed across a cluster
  - **Map**: Process chunks (splits) of data in parallel creating intermediate results as key-value pairs
  - **Shuffling**: Pairs with the same key are sent to the same reducer
  - **Reduce**: Process in parallel each of the pairs having the same key
- Hadoop: Apache's implementation of MapReduce
- YARN: Performs resource management and job scheduling for Hadoop
- Apache Spark: For in-memory data processing




# HDFS – Hadoop Distributed File System (Big Data Storage)

- Inputs and outputs of a map-reduce operation are stored in HDFS
- **Namenode** is a master node that contains metadata about how data is to be distributed in the cluster
- Here a single large file is split into five parts indicated by block-ids 1 to 5. Blocks 1 and 3 have two copies while 2, 4, 5 are replicated three times among the Datanodes
- Replication of data ensures fault-tolerance (default replication value is 3 in Hadoop)

## Block Replication

```
Namenode (Filename, numReplicas, block-ids, ...)  
/users/sameerp/data/part-0, r:2, {1,3}, ...  
/users/sameerp/data/part-1, r:3, {2,4,5}, ...
```

## Datanodes




[https://hadoop.apache.org/docs/r1.2.1/hdfs\\_design.html](https://hadoop.apache.org/docs/r1.2.1/hdfs_design.html)



# Spark and PySpark

- Spark accesses data from HDFS but avoids MapReduce processing, instead performing in-memory computing instead of on disk.
  - hence a faster alternative
- PySpark is a Spark API for Python
- A SparkContext represents the connection to a Spark cluster
- Subpackages:
  - pyspark.sql module
  - pyspark.streaming module
  - pyspark.ml package
  - pyspark.resource module



<https://spark.apache.org/docs/latest/cluster-overview.html>



# pyspark.sql

- This provides SQL-like functions for Spark DataFrames
  - SparkSession – create this to start working with a DataFrame
  - createDataFrame – via a variety of file formats such as csv, JSON, ORC, Parquet, Hive tables, Avro

```
data = [('John', '', 'Smith', '1993-03-06', 'M', 4000),  
        ('Gareth', 'Smart', '', '2002-07-29', 'M', 5000),  
        ('Linda', '', 'Carey', '1988-05-06', 'F', 5000),  
        ]
```

```
from pyspark.sql import SparkSession  
spark = SparkSession.builder.appName('Name of App').getOrCreate()  
columns =  
["firstname", "middlename", "lastname", "dob",  
 "gender", "salary"]  
df = spark.createDataFrame(data=data, schema =  
    columns)
```



# pyspark.sql

Note that Spark Dataframes are immutable, so any modification needs to be written to a new dataframe

#update a column

```
dfnew = df.withColumn("salary",col("salary")*100)
```

#show all rows

```
dfnew.show(truncate=False)
```

#use select for column subsetting, filter for row subsetting

```
dfnew.select("salary")
```

```
dfnew.filter(dfnew."salary" > 2000)
```





# Big Data – Example tools

- Hue, Apache Impala – real-time SQL for data warehouses (the latter for Hadoop)
- Pig/Hive: SQL-like scripting and querying tools for data processing that simplify MapReduce programs.
- HBase, MongoDB, Elasticsearch: Examples of a few NoSQL databases.
- Mahout, Spark ML: Tools for running scalable machine learning algorithms in a distributed fashion.
- Flume, Sqoop, Logstash: Data ingestion and integration (e.g. transfer of data between databases)
- Splunk, Elastic Stack: Monitoring of log files



# Amazon EMR

- Amazon Elastic Map Reduce – distributes processing over EC2 clusters
- Master node – manages the cluster
- Core nodes – run tasks and store results in HDFS
- Task nodes (optional) – run computation but do not make use of HDFS storage
- Spark Livy – a web service that connects with Spark over a REST interface enabling one to connect SageMaker to a cluster

Node type	Instance type	Instance count	Purchasing option
<b>Master</b> Master - 1	<b>m4.large</b> 2 vCore, 8 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	1 Instances	<input checked="" type="radio"/> On-demand <input type="radio"/> Spot Use on-demand as max price
<b>Core</b> Core - 2	<b>m4.large</b> 2 vCore, 8 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	2 Instances	<input checked="" type="radio"/> On-demand <input type="radio"/> Spot Use on-demand as max price
<b>Task</b> Task - 3	<b>m5.xlarge</b> 4 vCore, 16 GiB memory, EBS only storage EBS Storage: 64 GiB Add configuration settings	0 Instances	<input checked="" type="radio"/> On-demand <input type="radio"/> Spot Use on-demand as max price



# Lab 10.3: Big Data Computing

- Purpose
  - Learn how to use PySpark to analyse a large dataset
- Resources
  - A population dataset from Population Division, United Nations
- Materials
  - Jupyter Notebook (Lab 10\_3) in Google Colab



# Stream Computing

- Data Stream Processing
- The Pub/Sub Model
- Apache Kafka
- Amazon Kinesis
- Amazon Kinesis Analytics (SQL)
  - Windowed queries



# Data Streams are everywhere

- Sensor data
- Video
- News feeds
- Financial data
- Online gaming
- Network Traffic data
- Clickstreams or App activity tracking
- Social media
- Location tracking
- Machine logs



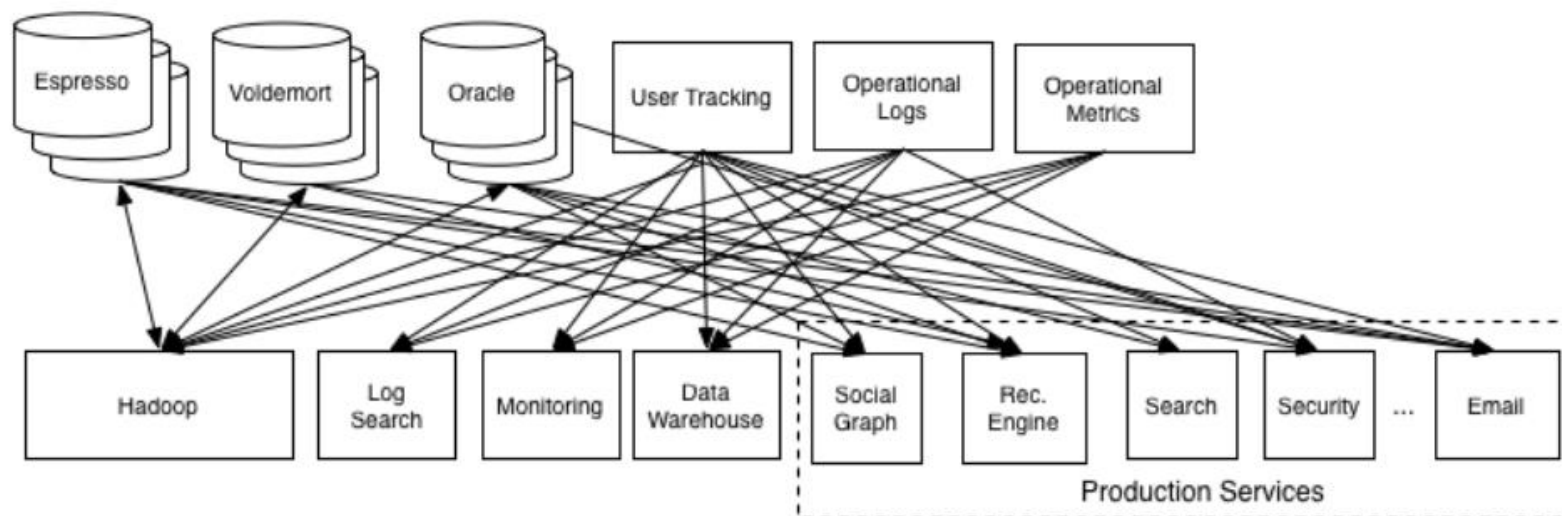
# Data Stream Processing

- Typically a solution comes from the use of multiple tools
- A stream process can be thought of as a directed acyclic graph composed of **sources**, **operators** and **sinks**
- **Sources** – where streams enter the streaming system
- **Operators** – functions of streams producing an output stream
- **Sinks** – places where streams flow out of the streaming system (e.g. a dashboard or database)
- Sources, sinks, operators combined into directed acyclic graph
- No batch processing, outputs constantly produced
- Event time and processing time generally differ
- Computation performed upon every event or within time windows



# The Pub/Sub Model

- Traditionally data systems and repositories are connected in the following manner, with a separate pipeline needed for each connection

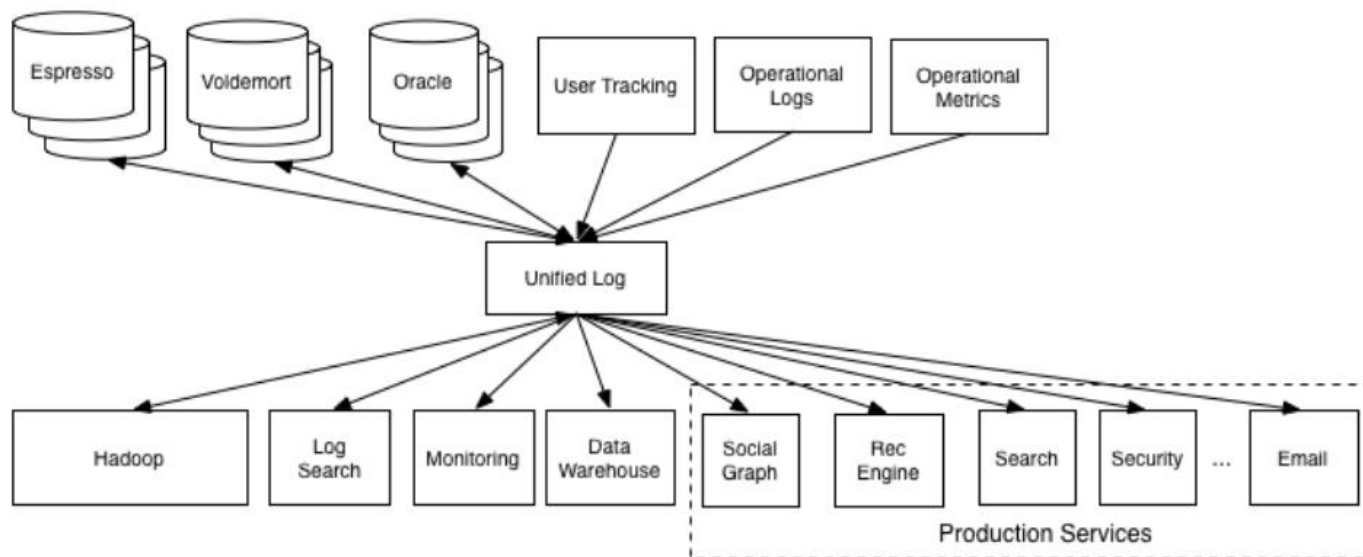


<https://engineering.linkedin.com/distributed-systems/log-what-every-software-engineer-should-know-about-real-time-datas-unifying>



# The Pub/Sub Model

- Connections can be greatly reduced by separating the **producers** and **consumers** of data, each publishing to/subscribing from a messaging system
- **Apache Kafka** is one such model employing distributed stream processing



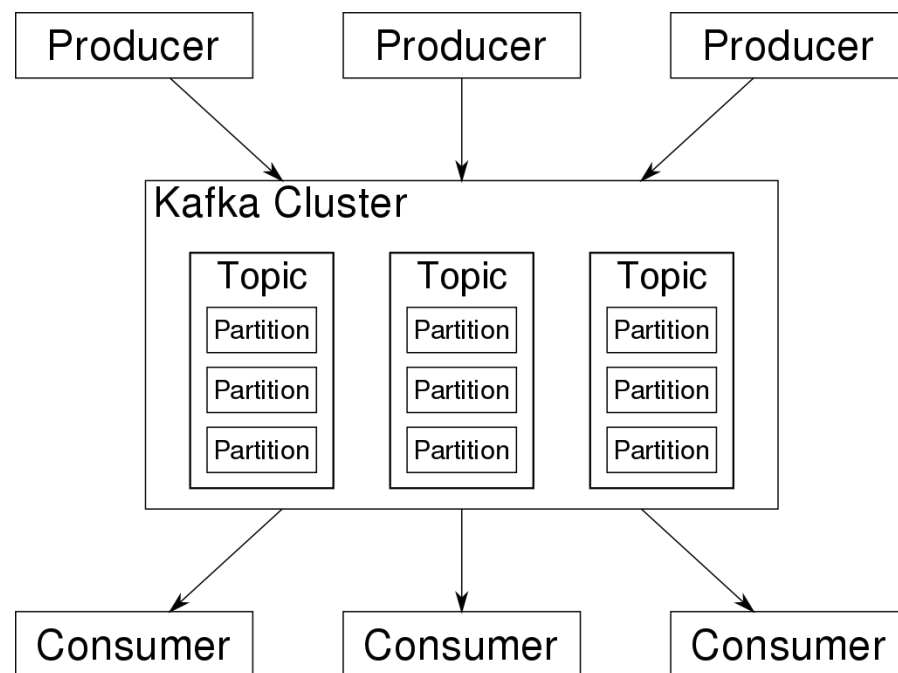
<https://engineering.linkedin.com/distributed-systems/log-what-every-software-engineer-should-know-about-real-time-datas-unifying>





# Apache Kafka

- **Broker** - handles all requests from clients (producers, consumers and metadata) and keeps data replicated within the cluster.
- **Topics** (categories) are divided into **partitions** containing records each with a unique offset in a fixed sequence. Partitions may be replicated across the cluster.






# Apache Kafka

- **Message (record)** - a key / value pair – the key assigns records to a specific partition (the publisher can also specify a partition directly)
- A record inside a partition has an **offset**, used to identify the order in which records are published to a partition – hence messages can either be read from the beginning or a specified offset
- **Zookeeper** – used for service synchronisation, tracking the status of nodes (brokers) in the Kafka cluster and maintaining a list of Kafka topics and messages (metadata)

## Anatomy of a Topic

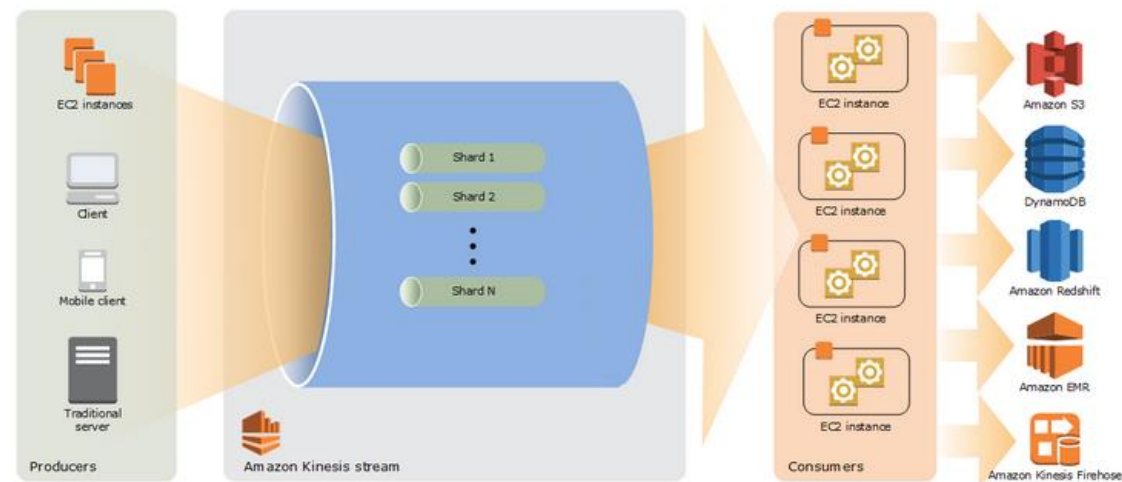


<https://sookocheff.com/post/kafka/kafka-in-a-nutshell/log-anatomy.png>



# Amazon Kinesis

- AWS platform for streaming data
- The throughput is in units of **shards**
- To put data into the stream, one specifies the **name** of the stream, a **partition key**, and the **data blob** to be added to the stream
- The partition key is used to determine which shard in the stream the data record is added to.



- 1 shard can support up to 5 transactions per second for reads, up to a maximum total data read rate of 2 MB per second and up to 1,000 records per second for writes, up to a maximum total data write rate of 1 MB per second (typically consumers are slower in transactions than producers)



# Lab 10.4 (Optional)

- Purpose
  - Learn how to use Apache Kafka in data streaming applications
- Resources
  - `fraud_data.csv`
- Materials
  - Jupyter Notebook (Lab 10\_4) in Google Colab



# Questions?



# Appendices



# End of Presentation!